

Functions & Functionals in R

PSS SUMMER SCHOOL

Anton Badev, Daniel Nikolic, Justin Skillman

June 22, 2017



Outline

- ▶ Loops
- ▶ Vectorization
- ▶ Functions
- ▶ Functionals
 - ▶ The `apply` family

Warm-up Exercise

- ▶ Take **5 minutes** to complete this task
- ▶ The Rdata file `m1_lecture2.Rdata` contains `ones_zeros` and `letts`
- ▶ They are both vectors of length 1,000
- ▶ Write a for loop to do the following:
 - ▶ If `ones_zeros[i]` is **1** change `letts[i]` into an upper case letter using `toupper()` function
 - ▶ If `ones_zeros[i]` is **0** leave `letts[i]` as a lower case
 - ▶ Record if a change was made
- ▶ Report $\frac{\text{number of changes made}}{\text{total length of letts}}$

Looping patterns

- There are three basic ways to loop over an object:

1. Loop over the **elements** of an object
2. Loop over **numeric indices**
3. Loop over the **names** of an object

```
> my_vec <- c(2,4,8)
> names(my_vec) <- c("first","squared","cubed")
> my_vec
  first squared cubed
     2       4      8
```

Looping patterns

Over the elements

```
for(x in my_vec) {  
  print(x)  
}  
[1] 2  
[1] 4  
[1] 8
```

Over numeric indices

```
n_elem <- length(my_vec)  
for(x in 1:n_elem) {  
  print(x)  
}  
[1] 1  
[1] 2  
[1] 3
```

Over the names

```
for(x in names(my_vec)) {  
  print(x)  
}  
[1] "first"  
[1] "squared"  
[1] "cubed"
```

Over the elements

- ▶ Can seem most natural to beginning programmers
- ▶ Becomes inefficient quickly
- ▶ *Example:* Take a uniformly random vector of integers and find the difference between each element and the last element in the vector

```
my_vec <- floor(runif(1000,1,100))
res_bin <- c() #empty vector for results
for(i in my_vec) {
  diff <- i - my_vec[length(my_vec)]
  res_bin <- c(res_bin,diff)
}
```

- ▶ R copies `res_bin` in memory every time you append to it

Over numeric indices

- ▶ Results are usually the same length as inputs because we want to perform some operation on *each* element
- ▶ Therefore we know the size of our output!
- ▶ *Example*: Take a uniformly random vector of integers and find the difference between each element and the last element in the vector

```
my_vec <- floor(runif(1000,1,100))
n_elem <- length(my_vec)
res_bin <- matrix(NA, n_elem, 1) #bin for results
for(i in 1:n_elem) {
  diff <- my_vec[i] - my_vec[n_elem]
  res_bin[i] <- diff
}
```

- ▶ Much faster, as R just assigns that index in `res_bin` a value

What is vectorization?

- ▶ Performs what would be a scalar operation on a vector
- ▶ Heuristic:
multiple statements e.g. loops → single statement
- ▶ Native to R environment
- ▶ Translates functions to C instead of R
- ▶ C takes advantage of BLAS (Basic Linear Algebra Subprogram)
 - ▶ SISD vs SIMD

Simple vectorization example

- Adding two matrices element wise

Scalar operations

```
a <- matrix(c(1,2,3,4),2,2)
b <- matrix(c(4,3,2,1),2,2)
y <- dim(a,1)
x <- dim(a,2)
res <- matrix(NA,x,y)

for (xi in 1:x) {
  for (xj in 1:y) {
    ans <- a[xi,xj] + b[xi,xj]
    res[xi,xj] <- ans
  }
}
```

Vectorized operation

```
a <- matrix(c(1,2,3,4),2,2)
b <- matrix(c(4,3,2,1),2,2)
res <- a + b
```

When to use vectorization?

- ▶ Great for replacing
 - ▶ Independent iterative processes
 - ▶ Making many repeated calls to functions in R (vs C)
- ▶ Not suited to replace
 - ▶ Dependent iterative processes
 - ▶ While loops
 - ▶ Functions with non-vector inputs
 - ▶ Algorithms where vectorization causes over-complication and readability issues

Advanced functions

- ▶ Components of functions
- ▶ Anonymous functions
- ▶ Closures
- ▶ Return values
- ▶ Editing a function

Components of a function

- ▶ `body()` → the function's code
- ▶ `formals()` → the function's arguments (inputs)
- ▶ `environment()` → the “sandbox” that maps the names of your function's objects to the locations of the objects themselves
- ▶ *Example:*

```
upp_or_low <- function(mtx, tri) {  
  if (tri=="upper") {  
    return(mtx[upper.tri(mtx)])  
  } else if (tri=="lower") {  
    return(mtx[lower.tri(mtx)])  
  } else {  
    stop("argument must be upper or lower!")  
  }  
}
```

Components of a function: example (cont.)

body

```
> body(upp_or_low)
{
  if (tri == "upper") {
    return(mtx[upper.tri(mtx)])
  }
  else if (tri == "lower") {
    return(mtx[lower.tri(mtx)])
  }
  else {
    stop("wrong arg!")
  }
}
```

formals

```
> formals(upp_or_low)
$mtx
```

```
$tri
```

environment

```
> environment(upp_or_low)
<environment: R_GlobalEnv>
```

Anonymous functions

- ▶ Functions that are not stored or given a name
- ▶ Used when a named function is not necessary (one time use)
- ▶ Often used in conjunction with functionals or closures
- ▶ *Example:* How can I assign `bin` the function's output and not the function itself using an anonymous function?

Which one will work? What will be their values?

```
>bin_1 <- function(x,y) (x + y)**3(2,1)  
>bin_2 <- (function(x,y) (x + y)**3)(2,1)
```

Anonymous functions

- ▶ Functions that are not stored or given a name
- ▶ Used when a named function is not necessary (one time use)
- ▶ Often used in conjunction with functionals or closures
- ▶ *Example:* How can I assign `bin` the function's output and not the function itself using an anonymous function?

Which one will work? What will be their values?

```
>bin_1 <- function(x,y) (x + y)**3(2,1)
>bin_2 <- (function(x,y) (x + y)**3)(2,1)

>bin_1(2,1)
Error in bin_1(2, 1) : attempt to apply non-function
>bin_2(2,1)
[1] 27
```

Closures

- ▶ A closure is a function that generates another function
- ▶ Anonymous functions are often *enclosed* in them
- ▶ Useful in conditional flow of a program
- ▶ *Example:*

```
>dffrnc <- function(offset) {  
  function(vec) {  
    start <- 1 + offset  
    end   <- length(vec) - offset  
    diff  <- vec[1:end] - vec[start:length(vec)]  
    return(diff)  
  }  
}  
  
>secdff <- dffrnc(2)  
>secdff(c(8, 10, 1, 7))  
[1] 7 3
```


Return values

- ▶ Return values can be specified using `return()` or if written as the last statement in the function
- ▶ `return()` is suggested in many style guides for readability

<pre>f <- function(z) { procedure_1 return(z) }</pre>	<p>equivalent</p> \longleftrightarrow	<pre>f <- function(z) { procedure_1 z }</pre>
--	---	--

- ▶ If you have more than one object to return you can do so in a `list()`
- ▶ Useful for returns of differing object type
- ▶ *Example:* For `my_optim` return the exposure matrix `exp_mtx` and eigenvector `eig_vec`

```
my_optim <- function(optim_args) {  
  optim_procedure  
  exp_mtx <- some_result  
  eig_vec <- some_other_result  
  return(list(exp_mtx, eig_vec))  
}
```

Editing a function

- ▶ The benefit of open source is that even if you didn't write it, you can modify the code
- ▶ Open source also means that *sometimes* there are errors or small bugs in a package
- ▶ Use `edit(<function>)` to edit a function in a text editor

Functionals

- ▶ **functionals** are functions that take another function as an input and return a “vector-like” object as output
- ▶ Example: `my_fnctnl <- function(g) g(1:10)`

```
my_fnctnl(mean)
[1] 5.5
my_fnctnl(sum)
[1] 55
```

- ▶ Commonly used in R for *vectorization*
- ▶ Allows for application of a function over data objects
 - ▶ `apply` family of functions

ifelse()

- ▶ `ifelse()` is a vectorized form of an `if else` statement
- ▶ `ifelse(boolean vector, function if T, function if F)`
- ▶ Example from warm-up:

```
ans <- ifelse(ones_zeros==1, toupper(letts), letts)
head(ans)
[1] "k" "S" "p" "E" "w" "w"
```

- ▶ Remember that this simply becomes a C `if` statement
- ▶ Avoid viewing functionals as a “black box”

The apply family

lapply

- ▶ *list apply*: list/vector input → list output
- ▶ C loop wrapped in R (much faster)
- ▶ Simple example

```
vec_list <- list(c(1:10),c(2:11),c(3:12))  
lapply(vec_list, function(x) max(x)**2)  
[[1]]  
[1] 100  
[[2]]  
[1] 121  
[[3]]  
[1] 144
```

- ▶ Each iteration uses a vector, `x`, as input

lapply (cont.)

- ▶ lapply with functions taking more than 1 input
- ▶ lapply **only** iterates over the list elements
- ▶ We can specify other arguments using the format `lapply(list, function, other_args)`

```
vec_list <- list(c(1:10),c(2:11),c(3:12))  
lapply(vec_list, function(x,pwr) max(x)**pwr, pwr=3)  
[[1]]  
[1] 1000  
[[2]]  
[1] 1331  
[[3]]  
[1] 1728
```

- ▶ Note: `unlist` coerces a list into a vector

The apply family

sapply

- ▶ Works the same way as lapply but returns a **vector**
- ▶ Vectors are often more useful as output

```
vec_list <- list(c(1:10),c(2:11),c(3:12))  
sapply(vec_list, function(x,pwr) max(x)**pwr, pwr=3)  
[1] 1000 1331 1728
```

The apply family

apply

- ▶ Applies a function over a dimension of a matrix
- ▶ `apply(m, dim, func, func_args)`
- ▶ Used often for summarizations or data manipulations

```
mtx <- matrix(c(1,2,3,4),2,2)
```

```
mtx
```

```
      [,1] [,2]  
[1,]     1     3  
[2,]     2     4
```

```
apply(mtx, 2, sum)
```

```
[1] 3 7
```


The apply family

tapply

- ▶ Applies a function over multiple groupings of data
- ▶ Groupings of data displayed as factors
- ▶ `tapply(vec, factors, func)`
- ▶ `tapply()` vs. `by()`

```
incomes <- c(40,52,150,61,103,41,55,33,90,55)
states <- c("PA","NY","NY","PA","CA","DC","CA","DC","TX","TX")
tapply(incomes, states, median)
  CA    DC    NY    PA    TX
79.0  37.0 101.0  50.5  72.5
```

In class exercise

- ▶ Complete the following problems in a group, we will ask each group to speak in **15 minutes**
- ▶ Babynames problem
 1. Load the package `babynames`
 2. Write a function that determines which name was the most popular in a given year. The only inputs to this function should be the vector of n 's for that year and the vector of `names`, the only output should be a character string that is the most popular name
 3. Write an `lapply` statement that finds the most popular name in all of the years in the dataset.
- ▶ Integer problem
 1. Write a function saved as `find_small_div` that takes in n , a , and b and returns the smallest integer in the set of integers between a and b that is divisible by n .
 2. Draw an integer vector of length 100 from a uniform distribution with support from 50 to 8,000. Save this as `samp_vec`.
 3. Take as given that a and b are 10,000 and 100,000 respectively and write an `apply` statement that applies `find_small_div` over `samp_vec`.

Summary

- ▶ Loops
- ▶ Vectorization
- ▶ Functions
- ▶ Functionals
 - ▶ The apply family